

# Week 6 - Pandas

(and a HINT of SQL...)

# Data Handling

What are the ways that we have learned so far to handle data?

- Lists of lists
- Dictionaries
- Custom Classes

None of these are particularly conducive to data exploration and quick manipulation

# Introducing Data Frames

When we want to manipulate data sets in a clean and efficient manner, we want to start thinking about data in terms of vectors:

- Each variable can be considered a vector
- Operations on a variable can be applied to all observations uniformly
- We can quickly reduce the number of variables for specific questions

*No way do we want to write a `for` loop for every little thing!*

# Introducing Data Frames

In Python, the `pandas` library contains the necessary code to begin working with Data Frames. It is dependent on many functions in the `numpy` library.

```
import pandas as pd # Import the library for use
```

# Creating a Data Frame

Create an empty Data Frame:

```
data = pd.DataFrame()
```

A Data Frame is a class that accepts the following parameters:

- `data`
- `index` (for referencing individual rows)
- `columns` (so you can name your variables)
- `dtype` (specify the **kind** of data for each column/variable)
- `copy` (whether or not the data should be duplicated in memory)

# Creating a Data Frame

We can also use pandas to easily read many types of files, and use them to create Data Frames containing the information stored in the file:

```
# CSV
data = pd.read_csv(your_filename_here.csv)
# or Excel
data = pd.read_excel(your_filename_here.xlsx)
# or STATA
data = pd.read_stata(your_filename_here.dta)
# or SAS
data = pd.read_sas(your_filename_here.sas7bdat)
# or SQL
data = pd.read_sql(your_query_here, your_connection_here)
# and many others!
```

# Referencing a Single Column

To access a list of all of the column names in your Data Frame:

```
data.columns
```

To then access (slice) a single column:

```
data['Column_Name']
```

To slice several columns at once into a new Data Frame, pass a list of column names:

```
data[['Column1', 'Column2']]
```

# Slicing the Data Frame - Hard Version

Two selection (or slicing) tools allow us to quickly subset our data.

```
data.iloc[row_selection, column_selection]
```

With the `.iloc` method, we can provide **integer**-based selections, or choose to select all rows or columns, and only subset on a single dimension.

```
data.iloc[:, 0] # Selects all rows, and first column
```



# Slicing the Data Frame - Easy Version

Two selection (or slicing) tools allow us to quickly subset our data across both axes.

```
data.loc[row_selection, column_selection]
```

With the `.loc` method (now with no `i`), we can provide **name**-based selections, choose to select all rows or columns, and more.

```
data.loc[:, 'ColumnName'] # Selects all rows, one column
```

# Slicing the Data Frame

With the `.loc` method (now with no `i`), we can also provide **logic**-based selections, creating subsets based on conditions.

```
data.loc[data['Column1'] == some_value, :]  
# Selects only the observations (rows) where the  
# condition is met
```

# Transforming our Data

We can quickly transform the data in a given column using the slicing techniques from above:

```
# Log the values of a variable, replace old column  
data.loc[:, 'Column1'] = np.log(data['Column1'])  
  
# Difference two variables – make new column!  
data['newColumn'] = data['Column1'] - data['Column2']
```

Because the variable doesn't exist yet, we don't use the .loc syntax here. Instead, we just create a new column by naming it! Super easy!

# Transforming our Data

We can choose an index from among our columns, instead of the arbitrary ascending numbers assigned by default:

```
data.set_index('transaction_id')
```

Or, we can establish a multi-level index by passing a list of columns:

```
data.set_index(['year', 'month', 'day'])
```

Remember! Indices **MUST** be unique values! In the case of a multi-level index, the combination of values from the multiple columns must be unique for each row to be a valid index.

# Datetimes and Date Processing

Processing Datetimes is also easy with built-in Pandas functionality:

```
data['myDate'] = pd.to_datetime(data['stringDateColumn'],  
    format = '%Y%m%d', # Need to indicate the correct  
    errors = 'ignore') #    format for your data!
```

We can also parse the data into separate columns afterward:

```
data['week'] = data['myDate'].dt.week  
data['day'] = data['myDate'].dt.day
```

# More Date Processing

A full list of the ways you can process dates is available at [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/timeseries.html#time-date-components](https://pandas.pydata.org/pandas-docs/stable/user_guide/timeseries.html#time-date-components).

# Cleaning Data

There are many operations that are not reasonable to perform with missing data. Any numeric transformation will fail to provide useful output where missing values exist.

```
# Resolve missing values in ALL columns at once
data.fillna(0, inplace = True)
# fills ALL missing values, overwrites original data

# Resolve missing values in single column
data['Column'].fillna(method='pad') # fill values forward
# We can use method 'backfill' to use the NEXT value,
# and fill backwards
```

# Map and Apply

`.map()` enables us to apply a function to each value item-by-item in a single column

`.apply()` provides the same functionality column-by-column or row-by-row (use the `axis=1` argument to specify that you want to operate on one row at a time)

- Each row is the input to the function, rather than a single value from one column



# Generating Summary Statistics

Using the `describe` function will create summary tables easily, and Pandas can even export them to csv for use in reports (this is true of ANY data frame in general, too!).

```
data.describe()
```

If we want the table presented similar to academic journal formats, we can add a few arguments:

```
data.describe().T[['count', 'mean', 'std', 'min', 'max']]  
# We need to transpose the data using .T before  
# we can select the descriptive stats we want to keep  
# Add a .to_csv('myfile.csv') to that line to save
```

# Other summary functions

Count unique values in a field:

```
data['column_of_interest'].nunique()
```

Show unique values in a field:

```
data['column_of_interest'].unique()
```

Show dimensions of a Data Frame:

```
data.shape
```

# Using SQL with Python (optional!)

In order to handle data on a large scale, we frequently rely on SQL databases. In this class, we can practice with MySQL.

Here is a link to analogous code for many other database types:

<http://docs.sqlalchemy.org/en/latest/core/engines.html>

Install MySQL connectors with the following command:

```
pip install sqlalchemy mysql-connector-python
```

# Using SQL with Python

The first thing we need to do is to establish a connection to our database:

```
from sqlalchemy import create_engine  
  
engineStr = 'mysql+mysqlconnector://viewer:'
```

We are using `mysql` via the `mysqlconnector` module. Next, we provide our `username:password`, which in this case is "viewer," with no password, so we do not enter text after the colon.

# Using SQL with Python

The first thing we need to do is to establish a connection to our database (will vary based on SQL flavor and server):

```
from sqlalchemy import create_engine

engineStr = 'mysql+mysqlconnector://*username*:~*password*'
engineStr += '@35.202.92.40:3306'
```

We need to direct the connection to our server, which is hosted at `dadata.cba.edu`, and can be reached through port `3306`.

# Using SQL with Python

The first thing we need to do is to establish a connection to our database:

```
from sqlalchemy import create_engine

# SQL flavor, user, password
engineStr = 'mysql+mysqlconnector://*username*:~*password*'
engineStr += '@35.202.92.40:3306' # Server Address
engineStr += '/nfl' # Database Name

engine = create_engine(engineStr) # Start the Engine
```

Last, we just need to include the database that we wish to access on the server. In this case, we can use `NFL`

# Retrieve SQL Data with Pandas

Our next step is to write a `SELECT` statement using SQL, and then to pass it to Pandas for retrieval.

```
select = """SELECT * FROM game WHERE seas=2019"""  
data = pd.read_sql(select, engine)
```

# Want to learn a bit about SQL queries?

Feel free to take a look at some slides about writing SQL query code:

<https://goo.gl/Lq2yC5>



# PandaSQL and Data Cleaning

We can actually use SQL to clean our data within Pandas by making use of the `pandasql` library.

Get started by using the following code:

```
from pandasql import sqldf
pysqldf = lambda q: sqldf(q, globals())
```

If it isn't installed, you can install the library by running

```
!pip install pandasql # "!" only needed in mimir/notebooks
```

# PandaSQL and Data Cleaning

```
edited_data = pysqldf(select_statement_here)
```

Using SQLite syntax, we can then clean any dataset using the same tools that we would to extract data from a database!

We can aggregate, create new columns, group, and join across datasets, just like we would with SQL.

**Lab Time!**