

Natural Language Processing with SpaCy

What even is it?

Natural Language Processing (NLP) is a tool similar to regex, but allowing us to take our text analytics to entirely new levels.

What even is it?

Regex	NLP
Create patterns to match in text	Identify the structure of text and use that to refine information
Used to verify or find data	Used to analyze data
Uses user-defined rules	Relies heavily on ML-based (or other) models

Tasks we can complete with NLP

- Identify parts of speech
- Iterate over sentences
- Find words used to describe various nouns (or anything else!)
- Visualize semantic relationships
- Filter text for analysis
- Conduct sentiment analysis

Getting Started

First, we will need to install the right libraries (and a corpus!):

```
# Load the SpaCy library

!pip install spacy
!python -m spacy download en_core_web_sm
```

A corpus is the body of knowledge that our NLP library (`spacy`) needs in order to be able to process the english language. We will use the smallest english corpus for now.

Get some text!

```
import requests

jane = requests.get(
    "https://github.com/dustywhite7/Econ8320/raw/master/AssignmentData/janeEyreCh1to3.txt"
).text
```

This file contains chapters 1 to 3 of Jane Eyre

(Lots of fun texts to play with at [Project Gutenberg](#))

Rev the NLP engines!

```
import spacy  
  
nlp = spacy.load("en_core_web_sm")  
doc = nlp(jane)
```

Here, we are pointing to the language file we want to use, and then parsing our text

Important attributes

Our new parsed document (`doc` in this case), has some important **attributes**:

- `sents` - a generator function to iterate over each sentence in the document
- `token` - each individual element of the document
 - Elements exist at the word/punctuation level

[Some] Attributes of tokens

- `lemma_` - the "root word" from which a token/word is derived
- `pos_` - the part of speech of a token/word
- `dep_` - the relationship of dependent tokens to the parent token (adjectives to nouns, etc.)
- `like_email` / `like_num` / `like_url` - check if a token is like an email, number, or url

Doing NLP

```
[i.lemma_ for i in doc if (not i.is_punct) and (not i.is_space)][:100]
```

```
['chapter',  
 'I',  
 'there',  
 'be',  
 'no',  
 'possibility',  
 'of',  
 'take',  
 'a',  
 'walk',  
 'that',  
 ...]
```

What is happening here??

Lemmatization

When we have words with many forms (verbs, plurals, etc.), we want to work with the **lemma** of that word, so that we can recognize that all forms are actually the same word being used in different ways.

"Ran" --> "run"

"Are" --> "be"

Mapping a sentence

```
from spacy import displacy

sent = [i for i in doc.sents][100]
displacy.serve(sent, style="dep")
```

Sentiment Analysis

Here, we are going to use `

Install `textblob`

```
! pip install textblob  
! python -m textblob.download_corpora
```

Note: only use the `!` characters if installing from inside of jupyter notebook or other python interpreter (not when installing from shell/command prompt)

Use `textblob` for sentiment analysis

```
from textblob import TextBlob  
  
blob = TextBlob(doc.text)  
  
for sentence in blob.sentences:  
    print(sentence.sentiment)
```

Lab Time!